

Conformity review v4

Sources as staged from the Mac at the time of writing. Compiles, 41 pages, body 6,600 words against 8,800, §8 still empty.

Two new defects, both introduced by today's edits

- N1.** `sections/eight.tex` lines 5 and 8 contain a literal `--`. The subsections *Comments on the Bayesian results* and *Consequences for the control agent* had an en-dash as their entire body, and both rendered. *Removed*; the two headings are now empty and waiting for §8.
- N2.** `autogluonChronosDatasets2024` is cited and not defined. It appears in `sections/retraining_appendix.tex` line 9 and has no entry in `patchAliasing.bib`, so it typesets as `[?]` on page 41. It is the only undefined reference in the log.

Verdicts

#	Requirement	Verdict	Where it stands
1	Operator, patch identity	MET	The false “unless $S = P$ ” clause is gone and the observability sentence now says, in §3, the sense in which the report uses the word. The thin part is deliberate: <i>indistinguishable</i> is answered by pointing at the representation level rather than derived.
2	f_s, P, S, f ; lock class	MET	Unchanged.
3	Semantic layer	MET	<code>pv:Forecast</code> , <code>forecastFrom</code> , <code>informsAnomaly</code> , the second line of <code>eq:semantic-path</code> and the <code>RampEvent</code> case give the brief's second clause a referent in the formalism. Nit: <code>introduction.tex</code> still calls it “an OWL2 ontology” with no serialisation.
4	Falsifiable hypotheses	PARTIAL	θ_{lock} now has a pre-registered rule, the table separates behavioural from representational, $M1$ is oriented, and the $H2$ caveat is in. One item left: the $H3b$ identification rule was not applied, so §5 still claims the design supplies the patch series.
5	Controlled experiments	MET	§7 and Appendix F carry the design and the training recipe.
6	Behavioural + representational	PENDING	Instrumented and designed; §8 is empty.
7	Bayesian analysis	PENDING	Specification consistent, filter disclosed, thresholds all pre-registered. Waiting on the numbers and the five LOO comparisons.
8	Mitigation	MET	Geometric argument intact; verdict corrected to inconclusive.
9	Agentic AI	MET	All four verbs, now with a forecast object behind them.
10	NIST AI RMF	PARTIAL	Threats, surfaces, consequences and residual exposure are there. GOVERN / MAP / MEASURE / MANAGE are still only a column of <code>tab:riskRegister</code> : no prose, no subcategory citations, GOVERN never discussed.

Blockers

- B1.** The executive summary is still Lorem ipsum (`sections/summary.tex`), on page 1.
- B2.** ~~The two `--` placeholders in `eight.tex`.~~ Removed. The two headings are now empty until §8 exists.
- B3.** The missing `autogluonChronosDatasets2024` bibliography entry.
- B4.** §8 is empty, so every number quoted in §9, §10 and `tab:riskRegister` cites a section that says nothing, and the register caption still claims each entry is anchored to a posterior reported there.
- B5.** When §8 is written: a verdict for $H3$ -location, θ_P reported, and `elpd_diff±dse` for the five LOO comparisons §7 promises — including $M1$'s second leg, which is what now decides $M1$. Also a convergence sentence you can keep: §7 currently says only a posterior passing the checks is carried into the three-way rule.

Cheap, if the words are there

Body is 6,600 and §8 will add roughly 700, so about 1,500 words remain. In order of value: two paragraphs restructuring §10 under the four RMF functions with subcategory citations; the $H3b$ identification rule in §5; the phase-design count in §7; the twenty-two geometries of the full set listed; and either finishing Appendix E's planned extension with the decomposition figures or cutting its three subsections and the TO DO comment.