

Conformity review v5

Compiles, 43 pages, body 6,641 words against 8,800, §8 empty. Fix blocks for points 4, 5, 6, 8, 9 and 10 are in `fix_5_to_10.tex`.

New since v4

- N1. The trim of `eight.tex` removed the geometric argument.** What survives is the clause “its geometric basis lies in the stride branch”. The argument itself — that reducing S raises the stride fundamental f_s/S while f_s/P is invariant, so overlap cannot reach the patch branch, instantiated on the one-minute telemetry — is gone. That argument *is* requirement 8’s “principled explanation of why the mitigation does or does not address the effect”, and it was the only part of $M1$ independent of a fit. Point 8 drops to **PARTIAL** until it comes back. Block P8b restores it in 110 words.
- N2. The “full set of twenty-two” is wrong, in the report and in the code comment.** `probe_lib.MODELS_ALL` holds 22 pairs; `EXCLUDE_S= {5, 15, 28}` removes *five*; `BAYES_MODELS`, which is what `collect.py` fits, holds 17. So the full set adds *two* geometries to the fifteen of `tab:hfModels` — `p8-s4` and `p16-s4` — not seven. Block P5b states it correctly and tabulates all 22 with their status.
- N3. A second Pagani attribution is missing.** Beyond the band tasks, the phase design also comes from that paper: `collect.py` draws phases with `p1.phases_Sf`, commented as “the S_f cap of Pagani et al. Eq. 6”. Neither D2 nor D3 cites it. Block P5a fixes the phase gap and the attribution together.

Verdicts

#	Requirement	Verdict	Where it stands
1	Operator, patch identity	MET	Unchanged from v4.
2	f_s , P , S , f ; lock class	MET	Unchanged.
3	Semantic layer	MET	<code>pv:Forecast</code> and the <code>RampEvent</code> case are in.
4	Falsifiable hypotheses	PARTIAL	Only $H3b$ ’s identification rule is missing → P4a.
5	Controlled experiments	PARTIAL	Two documentation gaps, both now with numbers: the phase design → P5a, and the full set → P5b. The second is a factual error, not an omission.
6	Behavioural + representational	PENDING	§8 is empty. Independently, Appendix F still carries three “Planned extension” subsections and a T0 D0 comment → P6a cuts them and saves 220 words.
7	Bayesian analysis	PENDING	Specification consistent, filter disclosed, all thresholds pre-registered. Waiting on the numbers and the five LOO comparisons.
8	Mitigation	PARTIAL	Regressed: see N1 → P8b. P8a supplies the confound paragraph for the empty results subsection, reframed now that a positive slope is the predicted direction.
9	Agentic AI	MET	Represent, monitor, explain, act are all present; the explain clause is still one sentence with no artefact → P9a is optional.
10	NIST AI RMF	PARTIAL	<code>ten.tex</code> contains no occurrence of GOVERN, MAP, MEASURE or MANAGE in prose — they exist only as a column of <code>tab:riskRegister</code> → P10a and P10b. This is the largest remaining gap on the Computer Security half.

Budget

Body is 6,641 with §8 empty; the old §8 was 701 words. Taking every block in `fix_5_to_10.tex`, including P6a’s saving, brings the body to 7,586 before §8, so about 8,290 with it — inside 8,800 but with roughly 500 words of margin rather than 1,500. If that is too tight, P9a (118) and P4a (199) are the two most droppable; P10a, P10b and P8b are not, since they carry requirement clauses that are currently unanswered.

Tables cost nothing, so `tab:fullSet` and the band-task table added yesterday are free.

Blockers

- B1.** The executive summary is still Lorem ipsum, on page 1.
- B2.** The `autogluonChronosDatasets2024` bibliography entry is still missing; it is the only undefined reference in the log and renders as [?].

- B3.** §8 is empty, so §9, §10 and `tab:riskRegister` cite a section that says nothing.
- B4.** When §8 is written: a verdict for $H3$ -location, θ_P reported, `elpd_diff` $\pm$ dse for the five LOO comparisons, and a convergence sentence that matches what you actually do.

On the $[2, 150]$ restriction

It is consistent with Deliverable 1, which pre-registered that “the analysis is conducted only in the frequency range over which valid forecasts are observed, which we expect to be a subset of $[2, 250]$ Hz”. §7 now states the scope and Appendix F keeps the full band, which is the right split: the seven band tasks partition 2–250, and no Bayesian model depends on that partition, since Model B uses the local $f_c \pm 1$ Hz task. The one thing still to declare is the criterion that returns 150, rather than the number alone.