

Deliverable 1

Structural Aliasing Analysis in Time-Series Forecasting Using Chronos

Matteo Boniotti, Gianluca Brignoli and Federico Sabbadini*

July 10, 2026

This investigation formalizes the phenomenon of *structural aliasing* within patch-based time-series foundation models, focusing specifically on Amazon’s Chronos-Bolt framework.[1, 2] The project focuses on smart-grid load balancing.

Domain Description

To ground this abstract architectural evaluation, the study utilizes a specific physical domain as an empirical reference layer. This reference setting models the solar energy produced by a photovoltaic array attached to an internal load, a local storage battery, and the national electricity grid. The explicit operational objective within this system is to evaluate production variations over short-term future horizons so that the system can dynamically modify the internal load, attach and detach the battery unit, with the aim of maintaining internal grid stability and enabling precise anomaly detection.

Within this infrastructure, a control AI agent is responsible for executing real-time tasks: dynamic load-balancing, battery dispatch, state-monitoring and active anomaly detection, and internal grid stability maintenance.

This analysis is directly intended for researchers, engineers and practitioners utilizing large language models within the Chronos family for time-series analysis, representation learning, and forecasting, as well as individuals working in the field of energy management which are interested in characterizing the fundamental limitations, structural boundaries, and latent capabilities of these architectures.

Purpose and Scope

This project formalizes and empirically validates the phenomenon of *structural aliasing* in patch-based time-series foundation models, using Amazon Chronos-Bolt-Tiny as the target architecture.

Chronos-Bolt-Tiny operates with a fixed patch size of $P = 16$, stride of $S = 16$ and context window of $T = 512$ tokens as default settings; its architecture comprises four encoder blocks followed by four decoder blocks terminating in a direct quantile projection head.[3]

This study aims to provide a comprehensive understanding of the relationships between raw signal characteristics and the latent representations learned by the model, systematically identifying the critical

*This work is licensed under the Creative Commons Attribution 4.0 International License. Concurrently, all associated source code, probing scripts, and software artifacts generated throughout the project lifecycle are released under the open-source MIT license.

Copyright for any third-party material included in this work remains with its respective owners and must be respected accordingly.

During the preparation of this work, the authors used AI-based tools, including GPT-5.5, Gemini 3.1 Pro, Gemini 3.5 Flash, Sonnet 4.6, and Opus 4.7, to support language refinement and improve clarity. All generated or suggested content was carefully reviewed, revised where necessary, and validated by the authors, who assume full responsibility for the final content of this work.

links between sampling frequency (f_s), patch size (P), and stride (S) in order to propose actionable solutions to mitigate the phenomenon.

The core objective is to challenge the baseline assumption that pre-trained temporal foundation models inherently preserve frequency attributes when processing Nyquist-compliant signals.

Mathematical Formulation

The project provides a formal domain description through OWL2 ontology representation of: the physical system (photovoltaic array, battery storage, national grid, internal load), the signal (frequency, amplitude, timespan, phase) and its patched representation (patch size, stride, sampling frequency). See Appendix A for a formal description of the ontology.

The investigation focuses on the formal definition of the phenomenon of structural aliasing, starting from the tokenization operator and describing phase-locking conditions. The project then formalizes the relationship between the signal frequency, patch size, stride, and sampling frequency, providing a mathematical framework to analyze the conditions under which structural aliasing occurs.

Let $\mathbf{x} = \{x[n]\}_{n \in \mathbb{Z}}$ be a discrete-time signal representing the input time series. We define the patch-based tokenization operator $\mathcal{T}_{P,S}$ with patch length $P \in \mathbb{N}^+$ and stride $S \in \mathbb{N}^+$ as a mapping that transforms the signal $\mathbf{x}$ into a sequence of vector-valued tokens $\mathbf{T} = \{\mathbf{t}_k\}_{k \in \mathbb{Z}}$, where each token $\mathbf{t}_k \in \mathbb{R}^P$. The k -th token is defined as

$$\mathbf{t}_k = \mathcal{T}_{P,S}(\mathbf{x})[k] = \begin{bmatrix} x[kS] \\ x[kS + 1] \\ \vdots \\ x[kS + P - 1] \end{bmatrix}. \quad (1)$$

Considering a continuous-time periodic signal sampled uniformly at sampling frequency f_s , yielding the discrete-time sequence $x[n]$. We isolate a fundamental harmonic component as

$$x[n] = \sin\left(2\pi \frac{f}{f_s} n + \phi\right), \quad (2)$$

where f denotes the signal frequency satisfying the Nyquist criterion $f < \frac{f_s}{2}$, and $\phi \in [0, 2\pi)$ is an arbitrary phase offset.

Phase-locking occurs when the structural phase profile of the signal repeats exactly under a translational shift equal to the stride S . Evaluating a stride translation yields

$$x[n + S] = \sin\left(2\pi \frac{f}{f_s} (n + S) + \phi\right) = \sin\left(2\pi \frac{f}{f_s} n + 2\pi \frac{fS}{f_s} + \phi\right). \quad (3)$$

For exact alignment between the signal and the tokenization operator, the phase accumulation over one stride must correspond to an integer number of 2π -periods

$$2\pi \frac{fS}{f_s} = 2\pi c, \quad c \in \mathbb{N}^+. \quad (4)$$

As a result, the signal exhibits exact stride-period translational invariance

$$x[n + S] = \sin\left(2\pi \frac{f}{f_s} n + 2\pi c + \phi\right) = \sin\left(2\pi \frac{f}{f_s} n + \phi\right) = x[n]. \quad (5)$$

Solving Equation (4) for f yields the phase-locking frequency class $\mathcal{F}_{\text{lock}}$, which is defined as the set of frequencies producing indistinguishable token sequences:

$$\mathcal{F}_{\text{lock}} = \left\{ f \in \mathbb{R}^+ \mid f = c \frac{f_s}{S}, c \in \mathbb{N}^+, f < \frac{f_s}{2} \right\}. \quad (6)$$

— DA SISTEMARE —
 Ipotesi H1, H2, H3
 — DA SISTEMARE —

Signal Generation and Dataset

The project utilizes a combination of synthetic and real-world datasets to evaluate the presence and impact of structural aliasing.

The investigation employs synthetic signals to establish a controlled environment for testing the model’s response to known frequency components. To do so elementary sinusoidal signals are generated with varying frequencies, amplitudes, and phases to systematically probe the model’s ability to capture and represent these characteristics. A second set of synthetic data is generated using a light implementation of TSMixup[1], which creates mixed-frequency sinusoidal signals to simulate various signal conditions.

If previous analyses show evidence of aliasing, the project may analyze signals produced with KernelSynth[1] to test whether the phenomenon persists in more complex signals. The project will then focus on the analysis of real-world data, through the use of the `Dataset-SolarTechLab.csv` dataset, which contains one-minute multivariate telemetry stream of a photovoltaic array.[4]

Methodology

At first the Chronos Bolt-tiny model available on the HuggingFace Hub is used to evaluate the presence of structural aliasing in the default configuration, with patch size $P = 16$ and stride $S = 16$. To evaluate the presence and impact of structural aliasing, the project will analyze a set of synthetic signals with varying frequency characteristics. At first the Chronos Bolt-tiny model is tested over simple sine waves with known frequency sistematically varied between 2 Hz and 256 Hz with 1 Hz increments, those tests are used to evaluate the presence of strcuctural aliasing.

Stride and patch size are then systematically varied according to the value reported in Table 1 to explore the explicit impact of patch geometry on aliasing effects. Since variations in the values of stride and patch size require a new training of the model we have selected a restricted set of parameters for which the model is re-trained.

Run ID	Patch Size (P)	Stride (S)
1	16	4
2	16	8
3	16	12
4	16	16
5	8	8
6	24	24

Table 1: *Patch and stride parameters used for model training.*

To obtain a comprehensive understanding of the impact of stride and patch size on the presence of structural aliasing, we have selected a set of reduced stride and patch size values, and a set of increased stride and patch size values, to evaluate the possibility of mitigating the phenomenon. We have decided to re-train the model also for the default values of patch and stride, with the aim of having a baseline for comparison

with the other configurations. Further possible training of the model with patch size fixed to 16 and stride 1 and/or 15 to evaluate the impact of small, and the most extreme, variation of the stride on the presence of structural aliasing and performance of the model.

We also considered the case of fixed patch size to default value 16 and variable stride to evaluate the impact of stride on the presence of structural aliasing.

The training procedure is based over the information provided by the official documentation about Chronos using 10M data produced with TSMixup and 1M data produced by KernelSynth.[5].

— DA SISTEMARE —

Possible description of the number of epochs and other parameters of the training procedure.

— DA SISTEMARE —

Models Evaluation

After model re-training a second set of synthetic signals is used to evaluate the presence of structural aliasing in the new models. This second set of synthetic signals includes:

- **Sine Wave Signals:** simple sine waves with known frequency sistematically varieted between 2 Hz and 256 Hz with 1 Hz increments, to assess the model’s ability to capture and represent periodic components across a range of frequencies.
- **Time Series Mixup Signals:** 100 signals generated by the light version of TSMixup procedure, mixing multiple sine waves of different frequencies; to evaluate the model’s capacity to disentangle and represent complex frequency interactions.

This set of 355 test signals is used for all the re-trained models, to evaluate the impact of stride and patch size on the model’s ability to capture and represent frequency information.

Optionally, if previous analyses show evidence of aliasing, the project may analyze test signals from a third set, including 100 signals generated using the KernelSynth tool. Also this set of test signals is used for all the re-trained models, to evaluate the impact of stride and patch size on the model’s ability to capture and represent frequency information in more complex signals.

During the analysis a SEED is set to ensure reproducibility of the results.

— DA SISTEMARE —

The project will focus on the analysis of the phenomenon, through MDL probing [6] after each decoder block to verify the presence, or absence, of frequency information in the hidden states.

Layer-resolved probing is executed sequentially across the intermediate hidden states of the encoder and decoder blocks, plus at the final output probability distribution level. By mapping representations across these specific layers, the project isolates where task-relevant signal attributes are preserved or compromised under architectural synchronization constraints.

Bayesian analysis.

— DA SISTEMARE —

Security Implication and Risk Assessment

Finally, the project will include a formal risk assessment grounded in the NIST AI Risk Management Framework (AI RMF) [7] to evaluate structural aliasing as a systematic vulnerability within cyber-physical energy systems. This analysis will outline plausible threat models and attack surfaces, possible mitigation strategies (including stride reduction), and an evaluation of whether the proposed mitigation strategies effectively reduce residual exposure for the automated control AI agent.

The risk assessment will be structured around the following key points: context and attack surface, threat modeling, impact, mitigation and residual exposure.

Significance and Impact

This subject is relevant because it explores an unexamined systemic vulnerability native to patch-based temporal foundation architectures. By systematically characterizing the conditions under which structural aliasing emerges, this work will provide critical insights into the fundamental limitations of these models, directly informing the design of more robust architectures and training paradigms.

For professionals operating within this domain, the analysis could expose how algorithmic data-deletion anomalies can propagate through automated distribution pipelines. By evaluating how hidden token-space representation collapse can deceive a load-balancing AI agent into executing misaligned control choices, this work could provide insights needed for high-reliability, fault-tolerant smart-grid deployments.

The impact of this project could extend beyond theoretical understanding, as it has direct implications for the reliability and safety of real-world applications that depend on accurate time-series forecasting and representation learning, particularly in high-stakes domains such as energy management.

Appendix A: OWL2 Ontology

The core ontology architecture bridges the gap between cyber-physical energy systems and deep learning signal processing domain knowledge. It categorizes the application setting into five main conceptual pillars rooted under the ontological hierarchy:

- **Physical Subsystem (pv:PhysicalComponent):** Represents the hardware infrastructure of the micro-grid, including the solar array, electrochemical storage, consumer loads, and the macro grid tie.
- **Cognitive Layer (pv:ControlAgent):** The artificial intelligence system responsible for autonomous grid optimization, load balancing, dispatch scheduling, and proactive vulnerability mitigation.
- **Signal Processing Layer (pv:Signal, pv:PatchedSignal):** Tracks the raw telemetry inputs (both time-domain features and frequency-domain components) along with their tokenized representation configurations (patch dimensions, stride values) tailored for the Chronos foundation architecture.
- **Environmental Tracking (pv:WeatherObservation):** Captures concurrent meteorological snapshots (air temperature, wind vectors) influencing structural physics and electrical efficiency.
- **Inferred Operational States:** Dynamic metadata classified into generation levels (pv:GenerationState), systemic disruptions (pv:AnomalyState), control directives (pv:AgentAction), and signal degradation hazards (pv:AliasingEvent).

Generation States

A formal description of the state of the photovoltaic array is provided through the following axioms, which define the generation state based on the measured tilted irradiance in W/m^2 :

$$\begin{aligned} \text{PhotovoltaicArray}(?pv) \wedge \text{measures}(?pv, ?s) \wedge \text{tiltedIrradianceWm2}(?s, ?g) \wedge (?g \geq 800.0) \\ \longrightarrow \text{hasGenerationState}(?pv, \text{ClearSky}) \end{aligned} \quad (7)$$

$$\begin{aligned} \text{PhotovoltaicArray}(?pv) \wedge \text{measures}(?pv, ?s) \wedge \text{tiltedIrradianceWm2}(?s, ?g) \\ \wedge (?g \geq 150.0) \wedge (?g < 800.0) \\ \longrightarrow \text{hasGenerationState}(?pv, \text{PartialCloud}) \end{aligned} \quad (8)$$

$$\begin{aligned} \text{PhotovoltaicArray}(?pv) \wedge \text{measures}(?pv, ?s) \wedge \text{tiltedIrradianceWm2}(?s, ?g) \\ \wedge (?g > 0.0) \wedge (?g < 150.0) \\ \longrightarrow \text{hasGenerationState}(?pv, \text{Overcast}) \end{aligned} \quad (9)$$

$$\begin{aligned} \text{PhotovoltaicArray}(?pv) \wedge \text{measures}(?pv, ?s) \wedge \text{tiltedIrradianceWm2}(?s, ?g) \wedge (?g \leq 0.0) \\ \longrightarrow \text{hasGenerationState}(?pv, \text{Nighttime}) \end{aligned} \quad (10)$$

Anomaly States

Formal description of the anomaly states which could be detected in the photovoltaic array, based on the measured irradiance and the measured power.

Aliasing Event

Formal description of the aliasing event, detected when the frequency of the signal is part of the frequency class $\mathcal{F}_{\text{lock}}$.

Agent Actions

Formal description of the action taken by the control agent, based on the current state and aliasing event.

References

- [1] Abdul Fatir Ansari et al. *Chronos: Learning the Language of Time Series*. Nov. 4, 2024. DOI: 10.48550/arXiv.2403.07815. arXiv: 2403.07815 [cs.LG]. URL: <http://arxiv.org/abs/2403.07815> (visited on 05/13/2026). Pre-published.
- [2] Abdul Fatir Ansari et al. *Fast and Accurate Zero-Shot Forecasting with Chronos-Bolt and AutoGluon / Artificial Intelligence*. Dec. 2, 2024. URL: <https://aws.amazon.com/blogs/machine-learning/fast-and-accurate-zero-shot-forecasting-with-chronos-bolt-and-autogluon/> (visited on 05/13/2026).
- [3] *ChronosBoltPipeline / Amazon-Science/Chronos-Forecasting*. DeepWiki. URL: <https://deepwiki.com/amazon-science/chronos-forecasting/2.4-chronosboltpipeline> (visited on 05/13/2026).
- [4] Sonia Leva et al. *Photovoltaic Power and Weather Parameters*. IEEE DataPort, Sept. 23, 2020. DOI: 10.21227/42V0-JZ14. URL: <https://ieee-dataport.org/open-access/photovoltaic-power-and-weather-parameters> (visited on 06/01/2026).
- [5] Abdul Fatir Ansari et al. *Chronos: Learning the Language of Time Series [SW]*. 2024. arXiv: 2403.07815. URL: <https://github.com/amazon-science/chronos-forecasting> (visited on 07/06/2026).
- [6] Elena Voita and Ivan Titov. “Information-Theoretic Probing with Minimum Description Length”. In: *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. EMNLP 2020. Ed. by Bonnie Webber et al. Online: Association for Computational Linguistics, Nov. 2020, pp. 183–196. DOI: 10.18653/v1/2020.emnlp-main.14. URL: <https://aclanthology.org/2020.emnlp-main.14/> (visited on 05/13/2026).
- [7] Elham Tabassi. *Artificial Intelligence Risk Management Framework (AI RMF 1.0)*. NIST AI 100-1. Gaithersburg, MD: National Institute of Standards and Technology (U.S.), Jan. 26, 2023, NIST AI 100-1. DOI: 10.6028/NIST.AI.100-1. URL: <http://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.100-1.pdf> (visited on 06/01/2026).