

Deliverable 2

Structural Aliasing Analysis in Time-Series Forecasting Using Chronos

Matteo Boniotti, Gianluca Brignoli and Federico Sabbadini*

August 17, 2026

This deliverable reports how the plan approved in Deliverable 1 was carried out, and it is written as a record of what was tested and how. Its thesis is that patch-based tokenisation makes certain frequencies poorly observable; that the affected set is fixed by the patch and the stride and splits into two branches the design keeps apart; and that the effect is measurable, partly mitigable by design, and representable in the semantics a control agent acts on. What follows runs from the trained models, to the frequency behaviour measured on them, to the Bayesian analysis that tests it, and then to the semantic layer and the security reading of the result.

Phase-Locking and Structural Aliasing

The two terms name different things, and separating them is what makes the hypotheses falsifiable rather than definitional. Phase-locking is a property of the input signal and the tokenisation geometry alone. It is exact arithmetic on (f, f_s, P, S) and holds for any model, trained or not. Deliverable 1 derives it as two conditions: $x[n + S] = x[n]$, under which consecutive patches are identical, and $x[n + P] = x[n]$, under which every patch holds a whole number of cycles. We call a frequency satisfying the first a stride lock and one satisfying the second a patch null, and the two sets they generate the stride branch and the patch branch of $\mathcal{F}_{lock}$. The set is their union, so a frequency may belong to one branch or to both. Structural aliasing is an empirical property of the learned representations, in which distinct inputs are mapped to nearly identical latent states and discriminative information is lost. Phase-locking makes it possible; it does not make it happen.

The split into two branches has a sharp experimental consequence. *H3* makes a separate claim about each branch, so evidence has to be attributable to one of them, and that requires frequencies where only one condition holds. When S divides P the stride branch is contained in the patch branch: every locked frequency then satisfies both, and no observation in that geometry can say which one produced a dip. Only $S \nmid P$ yields stride-only sites, and supplying them is what the design of Table 1 is built for.

Models

Table 1 lists the nineteen geometries the analysis runs on. All are trained from scratch under the same 100k-step budget and from the same seed, so that no comparison is confounded by training length, and all

*This work is licensed under the Creative Commons Attribution 4.0 International License. Concurrently, all associated source code, probing scripts, and software artifacts generated throughout the project lifecycle are released under the open-source MIT license.

Copyright for any third-party material included in this work remains with its respective owners and must be respected accordingly.

During the preparation of this work, the authors used AI-based tools, including GPT-5.5, Gemini 3.6 Flash, Sonnet 5, and Opus 4.6, to support language refinement and improve clarity. All generated or suggested content was carefully reviewed, revised where necessary, and validated by the authors, who assume full responsibility for the final content of this work.

P	S	O	tokens	$S \mid P$	stride-only
8	8	0.000	256	yes	0
16	8	0.500	255	yes	0
16	12	0.250	170	no	4
16	15	0.062	136	no	7
16	16	0.000	128	yes	0
24	8	0.667	254	yes	0
24	12	0.500	169	yes	0
24	15	0.375	135	no	6
24	16	0.333	127	no	4
24	20	0.167	102	no	8
24	24	0.000	85	yes	0
32	8	0.750	253	yes	0
32	12	0.625	169	no	4
32	15	0.531	135	no	7
32	16	0.500	127	yes	0
32	20	0.375	101	no	8
32	24	0.250	85	no	8
32	28 [†]	0.125	73	no	12
32	32	0.000	64	yes	0

Table 1: The nineteen geometries, all trained from scratch (seed 42) under the same 100k-step budget. $O = (P - S)/P$; the token count is at the 2048-sample training context and is what the overlap costs. The last column counts the in-band frequencies that are stride locks and not patch nulls: they exist only when $S \nmid P$, and the ten such runs supply 68 of them. [†] This stride does not divide the probing context: the run is swept but not fitted.

are published as one repository with a subfolder per geometry [1]. The repository also carries runs at the shortest strides; those are excluded here, matching Deliverable 1, because their stride-lock class has too few in-band members to carry a local $H1$ contrast. The published **chronos-bolt-tiny** checkpoint is not used at any point [2], so the $P=S=16$ baseline is budget-matched like the rest of the design.

The design is not symmetric for its own sake; each axis identifies something. Ten runs have a stride grid that is not contained in their patch grid, and between them they supply the 68 stride-only sites, the only frequencies at which a behavioural loss is attributable to the stride branch alone. Six strides appear at more than one patch size, which is what identifies the patch branch of $H3$ and makes δ_P estimable. The overlap ratio spans 0 to 0.75, and three runs share $O = 0.5$ at three different patch sizes, which is what separates an effect of the overlap ratio from an effect of the absolute stride. The token count in the fourth column is the price of that overlap: at a fixed context it grows as $1/S$, so the widest-overlap run carries four times the sequence length of the $P=S=32$ baseline.

One run is treated differently. The probing context has to close on a whole number of strides, or the last patch is internally padded and the padding, not the geometry, produces the collapse. That holds for every stride but the one marked with a dagger, whose smallest admissible context is far outside the probing convention: it is swept and reported descriptively, not fitted.

Table 2 writes the two branches out geometry by geometry: the spacings, the size of their union, and every in-band member of each. Bold marks a frequency both branches carry, which makes the identifiability argument readable by eye. A row whose stride column is entirely bold has no frequency belonging to the stride alone; the unbolded stride entries are the sites counted in Table 1.

Model	f_s/S	f_s/P	$ \mathcal{F}_{lock} $	stride branch	$c f_s/S$ (Hz)	patch branch	$k f_s/P$ (Hz)
p8-s8	64	64	3	64 128 192		64 128 192	
p16-s8	64	32	7	64 128 192		32 64 96 128 160 192 224	
p16-s12	42.67	32	11	42.67 85.33 128 170.67 213.33		32 64 96 128 160 192 224	
p16-s15	34.13	32	14	34.13 68.27 102.4 136.53 170.67 204.8 238.93		32 64 96 128 160 192 224	
p16-s16	32	32	7	32 64 96 128 160 192 224		32 64 96 128 160 192 224	
p24-s8	64	21.33	11	64 128 192		21.33 42.67 64 85.33 106.67 128 149.33 170.67 192 213.33 234.67	
p24-s12	42.67	21.33	11	42.67 85.33 128 170.67 213.33		21.33 42.67 64 85.33 106.67 128 149.33 170.67 192 213.33 234.67	
p24-s15	34.13	21.33	17	34.13 68.27 102.4 136.53 170.67 204.8 238.93		21.33 42.67 64 85.33 106.67 128 149.33 170.67 192 213.33 234.67	
p24-s16	32	21.33	15	32 64 96 128 160 192 224		21.33 42.67 64 85.33 106.67 128 149.33 170.67 192 213.33 234.67	
p24-s20	25.6	21.33	19	25.6 51.2 76.8 102.4 128 153.6 179.2 204.8 230.4		21.33 42.67 64 85.33 106.67 128 149.33 170.67 192 213.33 234.67	
p24-s24	21.33	21.33	11	21.33 42.67 64 85.33 106.67 128 149.33 170.67 192 213.33 234.67		21.33 42.67 64 85.33 106.67 128 149.33 170.67 192 213.33 234.67	
p32-s8	64	16	15	64 128 192		16 32 48 64 80 96 112 128 144 160 176 192 208 224 240	
p32-s12	42.67	16	19	42.67 85.33 128 170.67 213.33		16 32 48 64 80 96 112 128 144 160 176 192 208 224 240	
p32-s15	34.13	16	22	34.13 68.27 102.4 136.53 170.67 204.8 238.93		16 32 48 64 80 96 112 128 144 160 176 192 208 224 240	
p32-s16	32	16	15	32 64 96 128 160 192 224		16 32 48 64 80 96 112 128 144 160 176 192 208 224 240	
p32-s20	25.6	16	23	25.6 51.2 76.8 102.4 128 153.6 179.2 204.8 230.4		16 32 48 64 80 96 112 128 144 160 176 192 208 224 240	
p32-s24	21.33	16	23	21.33 42.67 64 85.33 106.67 128 149.33 170.67 192 213.33 234.67		16 32 48 64 80 96 112 128 144 160 176 192 208 224 240	
p32-s28	18.29	16	27	18.29 36.57 54.86 73.14 91.43 109.71 128 146.29 164.57 182.86 201.14 219.43 237.71		16 32 48 64 80 96 112 128 144 160 176 192 208 224 240	
p32-s32	16	16	15	16 32 48 64 80 96 112 128 144 160 176 192 208 224 240		16 32 48 64 80 96 112 128 144 160 176 192 208 224 240	

Table 2: The two lock branches of every geometry, in band $[2, 250]$ Hz at $f_s = 512$ Hz. The spacings f_s/S and f_s/P are in hertz and each branch is that spacing’s harmonics. $|\mathcal{F}_{lock}|$ counts their union, which is why it is not the sum of the two lists. **Bold** marks a frequency carried by both branches: when $S \mid P$ the stride branch is contained in the patch branch and every stride lock is bold, and only the remaining, unbolded stride entries are attributable to the stride alone.

Observed Frequency Behaviour

Every geometry is swept one hertz at a time over $[2, 250]$ Hz at $f_s = 512$ Hz, on a context of 480 samples with a 64-step horizon. The sweep grid is the uniform 1 Hz grid unioned with the exact predicted sites of every geometry, so no model is scored only where its own arithmetic predicts a dip and the non-integer sites are covered. Each swept frequency is measured three ways: the frequency the model regenerates, the amplitude it recovers relative to the injected one, and whether the input frequency can still be read back from the internal state by a probe. Locks are compared against two controls sharing the same background realisation and the same phase, placed at the largest offset not exceeding $0.25f_s/S$ that keeps both clear of either branch. As in Deliverable 1 the tone rides on 100 TSMixup and 100 KernelSynth signals normalised to unit variance, and the pure-sinusoid sweep is kept as the exact-degeneracy reference.

Lorem ipsum incididunt aliqua tenetur nostrud fuga labore saepe culpa do veniam quo optio. Est quis maiores cumque et doloribus ea deserunt et libero. Officia tempor alias dui deleniti optio irure sapiente ex. Dolores perferendis pariatur ullamco repellat sint quidem nostrud reiciendis dolore doloribus et proident rerum hic nihil molestias assumenda. Possimus dolorum corrupti at pariatur laborum ad corrupti itaque debitis officia placeat voluptatum aliqua ex voluptates optio cillum. Reprehenderit officiis optio eiusmod enim alias repellat molestias officia mollitia rerum possimus magna quis. Temporibus dolore et quas placeat praesentium expedita deserunt consectetur voluptas animi eu aliquip. Et qui voluptatum consequat at libero nam rerum ad

cillum placeat tempore reiciendis blanditiis aute impedit. Blanditiis eligendi laborum facilis anim reprehenderit minim nulla reprehenderit anim anim sit aut pariatur iusto praesentium ipsum. Optio hic harum doloribus excepturi commodo voluptates ut possimus perferendis nam. Tempore nam ullamco quibusdam tempore et sint magna sunt minus esse nisi officia ut ullamco. Asperiores reprehenderit tenetur exercitation fuga adipiscing aliqua sunt rerum. Vero deserunt fuga temporibus ea aliquip aut assumenda autem quibusdam quas. Irure ullamco officia iusto autem esse repudiandae consectetur proident earum. In delectus adipiscing earum corrupti quis accusamus repudiandae fuga cillum laborum deserunt hic sapiente. Provident deserunt occaecat est tempore mollit non repudiandae debitis laborum elit elit blanditiis repellendus accusamus occaecat deserunt. Mollitia fuga ad officia ullamco mollit repellendus cupidatat culpa proident quibusdam lorem autem deserunt minim ex. Non autem nulla quo similique veniam libero voluptas tempore minim velit eu commodo elit reprehenderit. Assumenda in temporibus mollitia voluptate reiciendis reiciendis consequat consectetur sit ullamco itaque aute quo occaecat culpa elit vero. Deleniti necessitatibus est cupiditate accusamus delectus optio consequat et animi omnis repudiandae. Necessitatibus consequat hic reprehenderit itaque eveniet amet minus pariatur ipsum reprehenderit.

Lorem ipsum ex alias et cupiditate repudiandae earum alias quibusdam laboris. Labore at sint ducimus eiusmod exercitation saepe facere perferendis elit dolore maxime cupiditate saepe voluptates non ducimus. Eveniet hic autem saepe at recusandae accusamus perferendis non placeat aute eligendi ea nam maxime molestias. Est nihil aliqua culpa quos ea voluptate fuga irure vero. Assumenda officia nostrud libero aut esse deserunt esse impedit voluptates cum. Optio cupidatat laborum excepturi quis fuga amet culpa maiores omnis minus amet expedita provident. Atque voluptates dolore nisi mollit ullamco minim iusto dignissimos eiusmod pariatur dignissimos consequat cumque accusamus cum reprehenderit. Voluptates repellat debitis cupiditate veniam blanditiis labore pariatur cumque aliqua odio amet veniam accusamus minim officia magna. Ea possimus dolor culpa maiores eligendi odio consequat tempor itaque est nisi esse. Incidunt pariatur non quas dolores earum proident deleniti placeat necessitatibus nulla dignissimos deserunt. Vero do sit amet saepe maiores sint voluptates temporibus. Placeat laboris impedit debitis delectus nam saepe dolores qui mollit officia cupidatat. Aute cum deserunt ut consequat sit aliqua eos impedit esse labore minim facilis saepe praesentium laborum atque tempor omnis excepteur. Odio placeat lorem iusto fuga provident reiciendis occaecati laborum sed quas. Laborum pariatur lorem similique facilis minim temporibus blanditiis necessitatibus non at saepe. Quis iusto veniam irure tempore eiusmod nam consectetur corrupti. Anim minim earum voluptate distinctio cupiditate debitis reprehenderit praesentium in tempor voluptates nihil. Saepe aute itaque saepe asperiores amet mollit minim elit eiusmod dui dolorum ullamco rerum facere alias incidunt amet hic laborum. Iusto lorem possimus magna necessitatibus tenetur quis itaque dolore temporibus vero enim iusto id proident anim. Omnis debitis facilis enim autem voluptatum tempor.

Lorem ipsum in provident eos quos asperiores dui sit quibusdam et aut odio exercitation qui officiis deleniti repudiandae. Voluptas assumenda assumenda ex reiciendis non quas minim temporibus amet deleniti omnis enim. Facere odio distinctio sunt sunt enim quis irure itaque iusto dolorum consequat eveniet blanditiis nisi fuga anim. Aut nam adipiscing velit lorem officiis facere cum quos irure eligendi deserunt rerum molestias ex provident. Cupiditate culpa libero ex cupidatat sit deleniti vero quidem. Nam distinctio enim dolorum nihil ducimus incidunt blanditiis ullamco ut. Voluptatum reprehenderit at odio quo eveniet molestias sint quidem nihil elit tempore maiores reiciendis proident ad incidunt nobis facere. Aute voluptatum aut incidunt reiciendis commodo eu repellendus eligendi officia praesentium corrupti eos accusamus cum est quos quibusdam. Nam ex cillum esse enim sunt necessitatibus rerum reiciendis officia facere similique facere nihil aute reiciendis occaecat. Quis fugiat officia alias quis excepturi est harum accusamus asperiores non consectetur. Nobis expedita nobis itaque sunt rerum dignissimos culpa et rerum blanditiis dolorum commodo necessitatibus earum qui quis dignissimos at expedita. Placeat impedit quas consectetur commodo sed cumque temporibus officiis lorem aliqua nam earum assumenda placeat. Laboris deserunt voluptate reprehenderit recusandae laboris omnis minim maiores eiusmod lorem commodo. Asperiores do quos commodo vero earum quo nisi exercitation aliqua corrupti itaque. Occaecat distinctio accusamus deserunt lorem dolor tenetur quos omnis blanditiis molestias laborum temporibus itaque id reiciendis at elit. Dolores labore consectetur occaecat

Model	Claim	Estimand	Predicts	Input
A	H1 behavioural	$e^{\bar{\beta}}$: recovery at a lock / at its controls	< 1	contrasts
B	H1 representational	$e^{\theta_{lock}}$: code length at a lock / elsewhere	> 1	mdl_cells
C	H2	σ_ϕ : spread of the deficit across phase	≈ 0	contrasts
D1	H3 location	θ_S, θ_P : dip depth on each grid, by LOO	both < 0	collapse
D2	H3a stride branch	κ_S : measured / stride-predicted spacing	$= 1$	sites
D2	H3b patch branch	κ_P : measured / patch-predicted spacing	$= 1$	sites
A'	M1 mitigation	δ_O : change in the deficit per unit of overlap	< 0	contrasts

Table 3: The models, the quantity whose posterior decides each claim, the value the claim predicts, and the probing table each is fitted to. H1 is tested twice, on the forecast (A) and on the representation (B), so that a discrepancy between them stays visible. H3 is split into its two branches, since $\mathcal{F}_{lock}$ is a union and a frequency may belong to either. M1 is not a separate model but the configuration level of A, promoted to a named estimand.

rerum optio ad eos mollit cumque harum mollit debitis sed culpa. Optio dolorum libero cupidatat ipsum deleniti.

Lorem ipsum proident debitis non quas occaecat anim assumenda officia iusto atque laboris debitis excepteur deserunt aut eligendi labore in. Ut culpa adipiscing irure eligendi ut et excepteur nam facere molestias nisi ad cillum provident. Excepteur itaque assumenda sed quas rerum quidem provident maxime eu laboris lorem. Blanditiis ad mollitia optio ea perferendis sunt facilis laborum quas. Veniam incididunt temporibus cupidatat quidem sapiente placeat occaecat occaecati fuga temporibus elit nobis at cum. Rerum sed voluptas dolore et eos occaecat dolore culpa. Dignissimos similique tempor iusto excepturi ducimus corrupti lorem dolore adipiscing anim laboris temporibus assumenda. Vero impedit debitis consequat rerum pariatur dolor quos reprehenderit id cupiditate excepturi omnis dolorum ad. Cupidatat nam velit at soluta dolore sed quibusdam maiores delectus cupiditate esse nihil ullamco aliqua iusto minim. Nostrud optio rerum placeat fugiat anim Duis eligendi omnis id tenetur ea. Atque blanditiis asperiores odio quidem eos accusamus cupidatat minus at excepteur laborum id. Praesentium sint cupiditate dolore libero vero laborum saepe itaque anim exercitation. Voluptas do ullamco ipsum temporibus anim placeat quidem eiusmod atque anim ex incididunt sint occaecat enim quidem voluptates nulla. Accusamus ipsum laboris mollitia qui do harum officia irure tempor proident eos do proident dolor cupiditate. Quidem excepteur quas enim proident sed debitis reiciendis quibusdam dolore soluta exercitation libero reiciendis voluptate. Hic quis esse libero dignissimos soluta praesentium dolores eligendi ut quas asperiores laborum eligendi.

[TO FILL, Collaborating author: the figure pages, from the current run of `reconstruction_figures.ipynb`. One full-width page per patch size, grouping that size's strides so that a feature moving between rows belongs to the stride branch; then the two cross-geometry summaries the notebook writes last. Each caption should state what the two panels plot, what changes between rows, and that the dashed lines are $\mathcal{F}_{lock}$ and the colour scale is probe accuracy from chance to one.]

Bayesian Analysis Design

Deliverable 1 applied a two-part framework to H1 alone. Here every claim is given its own measurement, its own model, and its own estimand: the single quantity whose posterior decides it. Table 3 lists the models with the value each claim predicts, Table 4 every prior, and Table 5 the rule that turns a posterior into a verdict. All are ordinary regressions; only the quantity on the left and the question the coefficient answers change between them.

Each model is built the same way: a measurement, a linear predictor μ_i that adds one offset per group the observation belongs to, and a likelihood. One coefficient is the estimand and the rest keep it unconfounded. Table 4 gives every prior with the role its parameter plays and the range over which the prior scale is varied as a sensitivity check; all effect priors are centred on zero effect, so no claim can be inherited from them. The

Model	Parameter	Prior	Role
A, C	$\bar{\beta}$	Student- $t_4(0, 0.5)$	population phase-lock effect, on the log-recovery scale
A	δ_O	Student- $t_4(0, 0.5)$	slope in $\widetilde{O}_c$, the centred overlap; one unit = 0.5 of O
A	δ_P	Student- $t_4(0, 0.5)$	slope in $\widetilde{\log P}_c$, the centred log patch size
A, C	$\tau, \sigma_k, \sigma_b, \sigma$	Half-Student- $t_4(0, 0.5)$	configuration, harmonic, background and residual scales
C	σ_ϕ	Half- $\mathcal{N}(0, 0.25)$	spread of the eight per-phase offsets
B	α_0	$\mathcal{N}(\log \bar{y}, 1)$	baseline log codelength
B	θ_{lock}	$\mathcal{N}(0, 0.5)$	log codelength expansion at a lock
B	k	Gamma(2, 0.1)	Gamma shape (dispersion)
B	$\sigma_{st}, \sigma_{geo}$	Half- $\mathcal{N}(0, 1)$	spread of the probe-stage and geometry offsets
D1	$\alpha_g, \theta_S, \theta_P$	$\mathcal{N}(0, 1)$	per-geometry off-grid level and the two grid labels
D1	σ	Half- $\mathcal{N}(0, 1)$	residual scale of $\log z_g(f)$
D2	κ_S, κ_P	$\mathcal{N}(0, 1)$	measured spacing per unit of predicted spacing
D2	σ_F	Half- $\mathcal{N}(0, 5)$ Hz	scatter beyond the sweep-resolution floor Δf

Table 4: Every parameter with its prior and its role. The effect priors are those of Deliverable 1, extended by δ_P , which the design of Table 1 makes estimable. All are centred on no effect, so a posterior that has not moved from its prior is itself a report of an absent effect. Prior scales are varied over $\{0.25, 0.5, 1.0\}$ as a sensitivity check.

fit returns a distribution rather than a point estimate, which is what permits the third verdict, inconclusive, that a pass/fail test cannot give.

H1

The behavioural measurement is the forecast amplitude recovery $R = A_{\text{pred}}/A_{\text{true}}$, collected in matched triplets: each lock f_k with its two controls, all three sharing the same background realisation and the same phase, which is what makes the contrast paired. The contrast is unchanged from Deliverable 1,

$$d = \log(R_k + 0.01) - \frac{1}{2} [\log(R_- + 0.01) + \log(R_+ + 0.01)], \quad (1)$$

one number per triplet, negative when the lock recovers less than its neighbours. It can be formed only because the injected tone has a known frequency. Model A fits it as

$$d_i \sim \text{Student-}t_4(\mu_i, \sigma), \quad \mu_i = \beta_{c[i]} + u_{k[i]} + u_{b[i]}, \quad \beta_c \sim \mathcal{N}\left(\bar{\beta} + \delta_O \widetilde{O}_c + \delta_P \widetilde{\log P}_c, \tau\right). \quad (2)$$

Three offsets build μ_i , one per configuration, lock harmonic and background draw, so that an unusual frequency or a single corpus cannot pass itself off as the effect. The heavy tails are the one deliberate departure from a Gaussian. Recovery falls to near zero where the model rebuilds nothing, and the logarithm then returns very large contrasts; under a Gaussian a few of them would set the answer alone.

The last term is the only unusual one, and it defines the estimand. Writing β_c as a draw from a distribution centred on $\bar{\beta}$ states that the geometries are instances of one phenomenon rather than unrelated experiments, and two things follow. First, $\bar{\beta}$ exists as a parameter: H1 is a claim about patch-based tokenisation, not about one checkpoint, so it needs a coefficient standing for the typical effect and carrying its own uncertainty. Second, each geometry weighs one unit. The number of usable triplets varies several times over with how

Claim	Rule	Prior	Reads as
H1 supported	$\Pr(\bar{\beta} < \log 0.8 \mid D) \geq 0.95$	0.34	at least 20% attenuation at a lock
H1 refuted	$\Pr(\bar{\beta} < \log 1.1 \mid D) \geq 0.95$	0.14	any effect is smaller than 10%
H2 supported	$\Pr(\sigma_\phi < \log 1.1 \mid D) \geq 0.95$	0.30	phase moves the deficit by under 10%
H3 location	two-label fit wins LOO, $\theta_S, \theta_P < 0$	—	dips sit on both predicted grids
H3a, H3b	$\Pr(\kappa_F - 1 < 0.1 \mid D) \geq 0.95$	0.05	that branch tracks its own parameter
M1 supported	$\Pr(\delta_O < 0 \mid D) \geq 0.95$, overlap term wins LOO	0.50	overlap reduces the deficit

Table 5: Decision rules, fixed before any posterior was inspected. **Prior** is the probability each rule already scored under the priors of Table 4: a posterior counts as evidence only in as much as it has moved away from that value. A claim is supported at 0.95, refuted at 0.05, and otherwise inconclusive. Support and refutation are separate rules, since failing to establish a 20% attenuation does not establish that there is none.

many lock sites fall in band, so pooling every row into one coefficient would let the largest design, not the largest effect, decide the answer. Nineteen configurations make that level worth fitting, and let δ_O and δ_P be estimated together instead of confounded. The overlap enters as a ratio and the patch size as $\log P$, because the patch grid has spacing f_s/P , so equal steps in $\log P$ are equal ratios of spacing. A tilde marks a covariate centred on the design, $\tilde{O}_c = (O_c - \bar{O})/0.5$ and $\widetilde{\log P}_c = \log P_c - \log \bar{P}$, which is what leaves $\bar{\beta}$ the effect at the average configuration instead of an extrapolation to $O = 0$ and $P = 1$.

Model B asks the same question of the representation, on a frequency-local prequential codelength: at each candidate f_c a probe must separate tones at $f_c \pm 1$ Hz from a captured internal state. Ten states are probed: the encoder [REG] token after each encoder block, the decoder state after each decoder block, the pre-projection vector and the quantile head. Chronos-Bolt carries a [REG] token in the encoder only, its decoder being driven by a single decoder-start token, so the decoder-side vectors are named decoder states. The band tasks of Deliverable 1 each summarise hundreds of frequencies at once, leaving the lock indicator nothing to attach to, and are kept as a cross-check. The regression is Gamma with a log link; $e^{\theta_{lock}}$ reads as how many times more bits a locked frequency costs, and probe stage and geometry enter as zero-mean offsets, since codelengths differ several-fold between stages.

H2

H2 uses the same triplets with one change. Equation 1 has already averaged over phase, so the phase index is kept instead and the response becomes the per-phase deficit $y_i = \log(R_{\text{ctrl},i} + 0.01) - \log(R_{\text{lock},i} + 0.01)$, which is d with its sign reversed and positive when the lock is worse. Model C is Model A on that response, without the configuration-level covariates and with the phase circle cut into eight slices, each given its own offset:

$$y_i \sim \text{Student-}t_4(\mu_i, \sigma), \quad \mu_i = \beta_{c[i]} + u_{k[i]} + u_{b[i]} + u_{p[i]}, \quad u_p \sim \mathcal{N}(0, \sigma_\phi). \quad (3)$$

The estimand σ_ϕ is the spread of those eight offsets, and it answers the question with one number. If the deficit depends on where in its cycle the signal starts, the offsets differ and σ_ϕ is large. If the geometry alone decides it, they are equal and σ_ϕ is near zero, as H2 predicts. Slices are used rather than a fitted curve because they assume nothing about the form a dependence might take. Dropping u_p gives a nested null, and comparing the two by leave-one-out separates the size of the phase effect from the evidence that it exists.

H3

H3 was stated in Deliverable 1 as two claims, one per branch. Both are tested on the token collapse $z_g(f)$, the dispersion of the input-patch embeddings across patches, which is zero when consecutive patches are identical. All geometries share one sweep grid: a uniform 1 Hz sweep unioned with the predicted sites of every

configuration. No geometry is therefore scored only where its own theory predicts a dip, and non-integer sites such as $42.\bar{6}$ Hz for $S=12$ are covered.

The first question is where the dips sit. Each swept frequency carries two labels, one per branch:

$$\log z_g(f) \sim \mathcal{N}(\alpha_g + \theta_S \cdot \text{OnStrideGrid} + \theta_P \cdot \text{OnPatchGrid}, \sigma). \quad (4)$$

The logarithm makes the labels multiplicative on the per-geometry off-grid level α_g . *H3* predicts both θ_S and θ_P negative, so the two-label fit should win the leave-one-out comparison against each label alone and against neither. The labels are not rivals for the same dips, since a frequency may carry either or both. But the two coefficients separate only where exactly one label applies, which for θ_S means the ten $S \nmid P$ geometries.

The second question is whether the dips move. A good fit above is still compatible with dips that happen to sit on a fixed grid, so each branch is checked against its own scaling law. Every detected dip is assigned to the branch that predicts it, and sites belonging to both are set aside. The fundamental of branch F is then compared with the spacing that branch predicts, $\Delta_g^S = f_s/S_g$ or $\Delta_g^P = f_s/P_g$:

$$\hat{f}_{1,g}^F \sim \mathcal{N}\left(\kappa_F \Delta_g^F, \sqrt{\sigma_F^2 + \Delta f^2}\right). \quad (5)$$

The floor Δf is the sweep step, since a spacing read off a discrete grid is no more precise than one step. Without it, sites landing exactly on the prediction drive σ_F to zero and the posterior degenerates.

H3a The stride branch. *H3a* predicts $\kappa_S = 1$: the stride sites track f_s/S one for one, so a site that stays put while S changes refutes it. It is identified by the series at fixed P , which the design supplies at every patch size in the table.

H3b The patch branch. *H3b* predicts $\kappa_P = 1$ in the same way. It is identified by the pairs that share a stride and differ in patch size, which the design supplies at six distinct strides.

Decisions and validation

Every threshold in Table 5 was fixed before a posterior was inspected, and validation comes before interpretation. Prior predictive checks precede collection; a parameter-recovery check then simulates data with a known effect and confirms the model finds it, which is what separates “no effect” from “a design that cannot see one”; posterior predictive checks are stratified by frequency, phase, P , S and generator. Competing formulations are compared by leave-one-out cross-validation rather than by null-hypothesis testing. Sampling uses NUTS with four chains, and a fit is reported only at $\hat{R} < 1.01$, effective sample sizes above 1000, and no divergences.

Mitigation Strategy

The mitigation carried through the design is increased patch overlap, that is a shorter S at fixed P . Its mechanism is arithmetic: the stride branch has its lowest in-band member at f_s/S , so halving the stride doubles that frequency and thins the family. The same arithmetic bounds what it can achieve. Overlap does not touch the patch branch at all, so it thins at best one of the two and leaves the P/k -minute blind periods of the photovoltaic mapping in place, and at a fixed context the token count grows as $1/S$, which is a real constraint for an edge-deployed forecaster. Both halves are testable rather than asserted, and $M1$ in Table 5 is how. Alternative mitigations, including stride randomisation, coprime-stride ensembling and phase-aware patch embeddings, are evaluated in the final deliverable.

Results

Lorem ipsum cupidatat nam cum proident ipsum quo velit cumque aliquip quis cum fuga omnis. Consequat sit ut maiores irure libero veniam repellat harum saepe eu. Mollitia praesentium esse recusandae eu magna laboris expedita officiis cupidatat quos. Tempor quibusdam molestias ut distinctio veniam esse officia cum cupidatat temporibus. Doloribus qui eiusmod tempore repudiandae velit expedita laborum ea reprehenderit at. Occaecat eiusmod perferendis do occaecati ex distinctio possimus reiciendis dolores optio dolores at cumque distinctio harum placeat necessitatibus minus nulla. Lorem officiis assumenda id placeat omnis nulla temporibus tempore. Magna commodo laborum impedit fuga quis maxime saepe eveniet eiusmod. Consequat minim molestias eveniet ad ut saepe rerum duis. Dolore nisi occaecat consequat officiis voluptatum cillum officia dolore. Vero velit occaecati ducimus possimus irure eos necessitatibus autem sunt iusto saepe id excepturi. Do cupidatat pariatur cum esse blanditiis cupiditate rerum eu iusto aliquip earum incididunt dolorum. Alias recusandae ullamco vero tenetur nam quidem iusto rerum harum in dolorum provident ad maxime mollit. Incidunt atque repudiandae vero quas molestias lorem sed officia reprehenderit deleniti. Impedit eligendi voluptates fuga incididunt consequat officiis mollit tempor consectetur ut lorem asperiores animi quos laboris recusandae laborum. Deserunt nobis quos duis proident fuga temporibus velit duis sit laborum reprehenderit facere nostrud dolore in dignissimos. Iusto dolor labore perferendis mollitia maxime repudiandae debitis at cillum lorem tempor et hic adipiscing. Excepteur id velit labore ullamco sit maiores cupidatat irure nobis non repudiandae saepe eligendi fugiat. Quas dolore corrupti incididunt autem tenetur ipsum rerum quo assumenda ad facere fugiat deserunt ullamco accusamus anim. Do ea similique iusto ut odio maiores quo recusandae iusto atque qui minim saepe sit eu accusamus.

Lorem ipsum cupiditate occaecat distinctio provident est facilis tenetur repellendus repellendus earum. Ipsum adipiscing quo anim repellat dolores culpa nam enim doloribus eu in sed adipiscing nisi laboris esse deserunt irure elit. Eiusmod aute eiusmod magna tempor dolore fuga non hic. Dolore expedita laboris at proident proident nisi sed sed veniam voluptatum autem exercitation consequat exercitation proident atque excepturi culpa. Accusamus consectetur mollitia eos praesentium incididunt harum occaecati necessitatibus temporibus voluptatum elit nobis elit quo. Exercitation deserunt repellendus incididunt tenetur doloribus qui quis voluptatum eu quo lorem itaque non voluptatum ut ipsum. Officiis nostrud officiis excepteur debitis deserunt voluptates accusamus velit praesentium culpa anim rerum cillum. Ad officiis perferendis ullamco cupiditate laborum nostrud tempore libero veniam. Adipiscing quidem proident quos iusto nihil delectus necessitatibus eu facilis anim omnis commodo hic sed. Cupiditate recusandae voluptate facere maiores occaecati eu voluptas minus eos anim commodo similique voluptas. Id saepe occaecat odio quos voluptate voluptate at cupiditate recusandae mollitia esse id cupiditate sint accusamus ullamco cillum ullamco. Expedita reprehenderit in quos corrupti quo ducimus cupidatat laboris laboris blanditiis proident. Voluptas sed sit tempore quo officia necessitatibus atque voluptas consectetur irure eos cum ipsum laborum. Repellat optio mollit mollit pariatur ea possimus impedit molestias accusamus exercitation optio laborum tempore velit. Cumque quibusdam possimus consectetur soluta repudiandae pariatur cupiditate dolor distinctio officiis laboris do. Delectus qui esse non repudiandae mollitia.

Lorem ipsum proident temporibus voluptates amet harum recusandae officia nobis. Possimus sunt excepteur nam voluptates ea laborum labore vero ducimus facilis tempore et sit enim optio optio animi iusto laboris. Quos tempore itaque officia nam voluptas culpa cillum consequat magna occaecat repellendus. Perferendis deserunt in animi nobis assumenda atque reiciendis commodo repellendus animi mollit odio rerum vero nihil excepteur quibusdam lorem. Blanditiis laborum laborum quos occaecati autem aut nihil minim dolorum voluptate quos expedita labore minim doloribus cupiditate aute earum deserunt. Sit dolor sunt aliqua atque vero exercitation irure anim excepteur facere deserunt voluptate sunt cum hic cillum quis reiciendis. Corrupti cupidatat debitis culpa earum ad minus aliquip alias ex iusto optio anim aute temporibus debitis alias labore quibusdam. Irure officiis at rerum cillum sapiente ipsum esse occaecati assumenda doloribus rerum atque assumenda quidem nihil. Enim pariatur dolorum elit consectetur tempor provident nostrud eveniet quibusdam aut irure sed culpa eligendi. Commodum culpa nostrud fuga officia temporibus itaque maiores sunt praesentium

quo officia cumque vero maiores ut deleniti deleniti animi. Cum similique necessitatibus dignissimos saepe deserunt proident debitis ex provident occaecat excepturi corrupti commodo veniam eiusmod. Maiores cum delectus repellat incididunt tempore corrupti laboris ipsum tempor sint temporibus et necessitatibus delectus. Rerum in minim culpa eiusmod omnis fugiat exercitation pariatur do optio exercitation sit harum aute quas preferendis accusamus. Excepteur optio sed excepturi consectetur impedit doloribus.

Lorem ipsum asperiores recusandae eiusmod ex optio cum placeat magna sit distinctio voluptate temporibus nobis reiciendis ullamco minim repellendus. Reprehenderit sit nihil ipsum dolor ea veniam qui ea consequat repellendus amet. Asperiores laborum facere excepteur incididunt fuga in minim atque alias rerum omnis eos. Sed dolor et sit ad distinctio quas quas cillum. Et molestias harum minus repellendus cillum in aliquip dolorum esse eligendi autem expedita facere dignissimos asperiores. Deleniti blanditiis et similique sit reprehenderit quas nihil at rerum distinctio rerum anim facere. Lorem occaecati iusto odio cumque velit eiusmod voluptatum irure repellat in ducimus reiciendis. Rerum deserunt hic minim sapiente maiores aut facilis non anim quas labore libero assumenda proident eos dolor expedita omnis. Veniam tenetur animi dolore anim libero recusandae accusamus recusandae occaecati autem saepe non sint culpa occaecat quis. Deleniti dolorum doloribus laborum cum repudiandae reprehenderit at tempor debitis quidem. Quidem voluptas ad voluptate molestias elit deserunt blanditiis repudiandae consectetur. Sed proident doloribus aut asperiores culpa accusamus blanditiis nihil nostrud. Consequat eos do culpa non pariatur rerum minim elit ut sed alias harum omnis aut dolore. Libero.

Semantic Layer

The ontology reproduced in Appendix A closes the loop from the signal to the decision. It represents the photovoltaic micro-grid, its telemetry and a single control agent, and it makes the tokenisation geometry a first-class object rather than a property of the model: a `pv:TimeSeriesSignal` has a `pv:PatchedRepresentation` carrying `pv:patchSize`, `pv:stride` and a derived `pv:overlapRatio`, and one signal may hold several, one per (P, S) under evaluation. Since the lock conditions are arithmetic on those same fields, the set of candidates is derivable from the knowledge base rather than hard-coded beside it, which is what lets an agent recompute it when a configuration or a sampling interval changes. Phase-locking is what the ontology infers; structural aliasing is what it records.

Where the Bayesian analysis enters is the point that matters here. An `pv:ObservabilityAssessment` does not read a state off the geometry. It records the estimand, its effect size, its credible interval and its posterior probability, and assigns `StructurallyAliased` only to a frequency that is a lock candidate and also carries confirmed evidence of loss, $A = C \wedge B$ in the notation of Appendix A. Geometry alone yields a candidate, never a verdict. Because the confirmation threshold exceeds one half, confirmed loss and confirmed preservation cannot both hold, and everything else falls to `PartiallyObservable`: an unfinished test, a credible interval crossing the threshold, or conflicting metrics. That is the inconclusive verdict of Table 5 carried into the semantics, and its consequence is operational rather than documentary, because the assessment selects the control mode. Confirmed aliasing forces the safe mode and suspends forecast-driven battery decisions; a merely partially observable assessment reduces the battery cap and lowers decision confidence. Uncertainty about the representation is therefore propagated to the actuator instead of being resolved by assumption.

Security Considerations

The forecaster is not a standalone artefact. It sits inside a control loop in which its output drives battery dispatch and curtailment on a photovoltaic micro-grid, so a forecasting failure is a control failure, and the property under test in this deliverable has to be read as a security property and not only as an accuracy one. The NIST AI Risk Management Framework [3] is the reference we adopt for that reading, because it separates the characteristics we can evidence here, validity and reliability, security and resilience, from the governance

functions that turn evidence into a control, and because it asks for known limitations to be mapped and measured rather than merely disclosed.

The limitation itself is unusual, and naming its shape correctly is what makes the rest follow. Structural aliasing is deterministic, attacker-independent and present at deployment time. It is not introduced by an adversary and it is not a defect in the training data; it is a consequence of the tokenisation geometry, and it is closed-form. The set of affected frequencies follows from (f_s, P, S) alone, and those three numbers are configuration, not weights. An adversary therefore needs no model access, no gradients, no query budget and no knowledge of the training corpus to know exactly where the model is blind. In the taxonomy of adversarial machine learning [4] this places the weakness with the availability violations rather than with integrity or privacy, and it places the attack in the white-box-free category, which is the part a conventional threat model tends to miss because it reasons about perturbations of the input tensor rather than of the physical process.

The attack surface follows directly. The perturbation an adversary needs is physical and small: a periodic component injected at, or dithered onto, a locked frequency by any controllable load, inverter switching pattern or curtailment schedule is present in the plant and absent from the forecast, because the tokenisation maps it to the same latent state as its absence. Two consequences matter operationally. The signal is invisible rather than misread, so residual-based anomaly detection downstream of the forecaster does not see it either; and the same arithmetic that makes the exploit cheap makes it stealthy, since the injected component is legitimate plant behaviour at a frequency the operator has no reason to treat as special. Detection has to be placed before tokenisation, on the raw telemetry, or it is placed inside the blind spot it is meant to guard.

What the framework then requires is that this be mapped, measured and managed rather than noted. Mapping is Table 1 together with the lock arithmetic, which enumerates the affected set for every deployed configuration. Measurement is the Bayesian analysis of Bayesian Analysis Design: the estimand and its credible interval are the evidence artefact, and the prior probability recorded beside each threshold is what makes the verdict auditable rather than asserted. Management has two halves, one in the model and one in the system. In the model it is the overlap mitigation, whose bound is stated with it, since overlap thins one branch and leaves the other untouched, so residual risk survives by construction and has to be reported as such. In the system it is the semantic layer, where a `PartiallyObservable` assessment degrades actuation instead of failing silently, which is the resilience property the framework asks for. The full risk assessment, with the threat model, the residual-risk register and the operational controls, is carried in the final deliverable.

Appendix A: Ontology of the Photovoltaic Application Domain

This document defines a compact ontology for a photovoltaic system governed by one control agent. It connects environmental and electrical observations to anomaly recognition, Chronos-style patched representations, structural-aliasing evidence, and bounded power-routing decisions [5]. Routing is included only as the operational endpoint of the semantic model; the principal concern remains whether the representation preserves the information on which the agent relies.

Domain Description

The modeled system contains four energy assets: a photovoltaic array, battery storage, an aggregate internal load, and an external grid. The internal load is a passive sink representing a factory or a fleet of machines. Its demand is observed, but it has no agent, machine-level taxonomy, curtailment command, or scheduling model. The external grid is an available bidirectional boundary bus. It supplies residual demand and accepts residual production.

Exactly one `pv:ControlAgent` coordinates the four passive assets. It observes one synchronized control cycle, assigns environmental, anomaly, and observability states, selects a control mode, and emits one or more directional decisions. There is no peer-to-peer negotiation among assets.

The seven admissible transfer directions are

$$\begin{aligned} & \text{PV} \rightarrow \text{Load}, & \text{PV} \rightarrow \text{Battery}, & \text{PV} \rightarrow \text{Grid}, \\ & \text{Battery} \rightarrow \text{Load}, & \text{Battery} \rightarrow \text{Grid}, & \text{Grid} \rightarrow \text{Load}, \\ & & & \text{Grid} \rightarrow \text{Battery}. \end{aligned} \tag{6}$$

They satisfy the cycle-level balance

$$P_{\text{pv}} + P_{\text{bat,dis}} + P_{\text{grid,in}} = P_{\text{load}} + P_{\text{bat,ch}} + P_{\text{grid,out}}. \tag{7}$$

All terms denote non-negative magnitudes; import and export are therefore kept separate rather than encoded by an ambiguous sign.

Purpose and Scope

The ontology provides the smallest semantic layer needed to represent the physical boundary, validated observations, environmental context, three selected anomalies, and the effect of patching on model observability. It does not model individual machines, market prices, dispatch optimisation, grid faults, or a deployed real-time controller. Load demand, battery SoC, asset limits, and the optional grid-exchange reference are asserted or simulated because they are not present in the photovoltaic telemetry dataset.

The environmental states are operational generation proxies, not meteorological diagnoses. Likewise, the routing rules describe a deterministic conceptual policy rather than the central experimental contribution. Structural aliasing remains an empirically tested property of a signal-representation pair; no fixed synthetic frequency is transferred to the one-minute photovoltaic signal.

Ontology Representation

Core classes and closed vocabularies Table 6 defines the compact class interface. The four asset subclasses form a disjoint union; each `pv:PVSys` contains exactly one asset of each type and is managed by exactly one control agent. Closed operational vocabularies are represented as named individuals rather than deeper class hierarchies.

Class	Role
<code>pv:PVSys</code>	Aggregate system and scope of each control cycle.
<code>pv:EnergyAsset</code>	Parent of the four passive physical assets.
<code>pv:PhotovoltaicArray</code>	Intermittent renewable source measured through PV telemetry.
<code>pv:BatteryStorage</code>	Bounded storage asset that can charge, discharge, or hold.
<code>pv:InternalLoad</code>	Aggregate passive demand of a factory or machine fleet.
<code>pv:ExternalGrid</code>	Reliable bidirectional boundary bus for residual balancing.
<code>pv:ControlAgent</code>	The single decision-making entity.
<code>pv:SystemObservation</code>	Time-stamped snapshot of PV, environment, battery, load, and grid reference.
<code>pv:TimeSeriesSignal</code>	Physical signal and its temporal descriptors.
<code>pv:PatchedRepresentation</code>	A model input representation defined by patch and stride.
<code>pv:ObservabilityAssessment</code>	Evidence linking a signal-representation pair to an observability state.
<code>pv:GenerationContext</code>	Operational proxy for expected solar generation.
<code>pv:AnomalyState</code>	Detected physical or sensing abnormality.
<code>pv:ObservabilityState</code>	Semantic assessment of information retained after patching.
<code>pv:ControlMode</code>	Normal, conservative, or safe operating policy.
<code>pv:PowerRoute</code>	Closed vocabulary of the seven admissible source–destination pairs.
<code>pv:PowerRoutingDecision</code>	One directional transfer or hold decision with an optional bounded setpoint.

Table 6: Core classes in the simplified ontology.

The closed vocabularies are:

Generation context: `ClearSky`, `PartialCloud`, `Overcast`, and `Nighttime`.

Anomaly: `InverterTrip`, `RampEvent`, and `SensorMalfunction`.

Observability: `Observable`, `PartiallyObservable`, and `StructurallyAliased`.

Control mode: `NormalMode`, `ConservativeMode`, and `SafeMode`.

Power route: `PVToInternalLoad`, `PVToBattery`, and `PVToExternalGrid`;
`BatteryToInternalLoad` and `BatteryToExternalGrid`;
`ExternalGridToInternalLoad` and `ExternalGridToBattery`.

Formally, each vocabulary class is an OWL-style enumeration of the listed named individuals (`oneOf`). The five vocabulary classes are pairwise disjoint, and all listed values are declared globally different (`AllDifferent`). This prevents a value such as `SafeMode` from being silently identified with `ClearSky` or `NormalMode` under open-world semantics. Each valid time-stamped observation has exactly one generation context. Each completed observability assessment has exactly one observability state, and each routing decision records exactly one control mode. Co-issued decisions based on the same synchronized observation must record the same mode. Anomalies are not declared mutually exclusive because, for example, an abrupt inverter trip also creates a large power change. Deterministic control is obtained through the priority rules in the closed-loop section.

Object properties The semantic relations in Table 7 connect the physical boundary to the inference and control layers. Each active `pv:PowerRoutingDecision` has exactly one route value, one source, and one distinct destination. Its non-negative setpoint magnitude is optional because `Balance` may denote passive residual exchange. A cycle that simultaneously serves the load, charges the battery, and exports a residual uses three instances of the same decision class, all based on the same observation and mode.

Property	Domain	Range or meaning
<code>pv:hasAsset</code>	<code>pv:PVSystem</code>	<code>pv:EnergyAsset</code>
<code>pv:managedBy</code>	<code>pv:PVSystem</code>	exactly one <code>pv:ControlAgent</code>
<code>pv:hasObservation</code>	<code>pv:PVSystem</code>	<code>pv:SystemObservation</code>
<code>pv:hasSignal</code>	<code>pv:PhotovoltaicArray</code>	<code>pv:TimeSeriesSignal</code>
<code>pv:hasRepresentation</code>	<code>pv:TimeSeriesSignal</code>	<code>pv:PatchedRepresentation</code>
<code>pv:assessesRepresentation</code>	<code>pv:ObservabilityAssessment</code>	exactly one <code>pv:PatchedRepresentation</code>
<code>pv:hasGenerationContext</code>	<code>pv:SystemObservation</code>	zero or one <code>pv:GenerationContext</code> value
<code>pv:hasAnomaly</code>	<code>pv:SystemObservation</code>	zero or more <code>pv:AnomalyState</code> values
<code>pv:hasObservabilityState</code>	<code>pv:ObservabilityAssessment</code>	one <code>pv:ObservabilityState</code> value
<code>pv:selectsMode</code>	<code>pv:PowerRoutingDecision</code>	exactly one <code>pv:ControlMode</code> value
<code>pv:makesDecision</code>	<code>pv:ControlAgent</code>	<code>pv:PowerRoutingDecision</code>
<code>pv:basedOn</code>	<code>pv:PowerRoutingDecision</code>	observation, assessment, or inferred state
<code>pv:hasRoute</code>	<code>pv:PowerRoutingDecision</code>	zero or one <code>pv:PowerRoute</code> value
<code>pv:hasSourceAsset</code>	<code>pv:PowerRoutingDecision</code>	zero or one <code>pv:EnergyAsset</code>
<code>pv:hasDestinationAsset</code>	<code>pv:PowerRoutingDecision</code>	zero or one <code>pv:EnergyAsset</code>

Table 7: Object-property signatures.

`pv:hasRepresentation` is intentionally non-functional because one signal can be evaluated under several (P, S) configurations. Each assessment identifies exactly which representation it evaluates. Each decision is based on exactly one current observation and its completed assessment; its own `pv:selectsMode` value therefore scopes the mode without accumulating incompatible modes on the persistent agent. The range of `pv:basedOn` is the union of observation and inferred-assessment entities, rather than the intersection that multiple RDFS range declarations would imply.

OWL’s open-world semantics do not by themselves validate that endpoint types match a named route. The operational validator therefore admits only the seven ordered pairs in Equation 6, rejects identical endpoints, and checks consistency among `pv:hasRoute`, source, and destination. A hold decision has no route or endpoints and has a zero or absent setpoint. The `pv:basedOn` relation is retained only as the minimum evidence link required to justify a selected mode; no separate trace model is introduced.

Data properties and units The synchronized observation combines the measured photovoltaic and environmental variables with asserted or simulated battery, load, and grid-reference values. The control-critical quantities are PV power P_{pv} , tilted irradiance G_{tilt} , normalized battery state of charge $s \in [0, 1]$, and aggregate load demand $P_{\text{load}} \geq 0$. Each has a separate validity flag derived from presence, finiteness, freshness, and a configured physical range. Horizontal irradiance, air temperature, wind speed, and wind direction remain environmental evidence but do not independently block the controller. The signed `pv:gridExchangeReferenceW` is optional, with positive values denoting export; it is usable only when present, finite, and fresh.

Missing or stale data cannot be inferred safely from OWL’s open-world assumption alone. An operational validation layer therefore compares the observation time stamp with a configurable staleness limit, validates numeric ranges, and asserts the four critical Boolean flags consumed by the ontology. This keeps data-quality validation distinct from semantic classification.

Owner	Data properties
pv:SystemObservation	pv:pvPowerW, pv:powerRampRateWPerS, pv:loadDemandW, optional pv:gridExchangeReferenceW, pv:tiltedIrradianceWm2, pv:horizontalIrradianceWm2, pv:airTemperatureC, pv:windSpeedMs, pv:windDirectionDeg, pv:batterySoC, pv:observationTimestamp, pv:pvPowerValid, pv:tiltedIrradianceValid, pv:batterySoCValid, and pv:loadDemandValid.
pv:BatteryStorage	pv:capacityWh, pv:reserveSoC, pv:targetSoC, pv:maximumSoC, pv:maximumChargePowerW, and pv:maximumDischargePowerW.
pv:TimeSeriesSignal	pv:signalFrequencyHz, pv:samplingFrequencyHz, pv:periodSamples, pv:amplitudeW, pv:timeSpanSeconds, and pv:phaseOffsetSamples.
pv:PatchedRepresentation	pv:patchSize, pv:stride, and derived pv:overlapRatio.
pv:ObservabilityAssessment	pv:analysisStatus, pv:assessmentMethod, pv:assessmentMetric, pv:strideLockingRatio, pv:patchLockingRatio, pv:lockingTolerance, pv:effectEstimate, pv:posteriorLossProbability, pv:credibleIntervalLower, pv:credibleIntervalUpper, and pv:assessmentTimestamp.
pv:PowerRoutingDecision	optional non-negative pv:powerSetpointW, pv:decisionConfidence, pv:batteryCommand in {Charge, Discharge, Hold}, pv:pvCommand in {Connect, Isolate}, and pv:gridCommand in {Import, Export, Balance}.

Table 8: Quantitative property groups and units encoded by their names.

A valid patched representation satisfies $P \geq 1$ and $1 \leq S \leq P$. Overlap is descriptive metadata:

$$\gamma = \frac{P - S}{P}. \quad (8)$$

It is not used as a standalone rule for declaring aliasing.

Inference Rules

Environmental generation context Let $g_0 < g_1 < g_2$ be site-configurable tilted-irradiance thresholds and let G_t denote the validated value of G_{tilt} at time t . The active generation context is

$$C_G(G_t) = \begin{cases} \text{Nighttime,} & G_t \leq g_0, \\ \text{Overcast,} & g_0 < G_t < g_1, \\ \text{PartialCloud,} & g_1 \leq G_t < g_2, \\ \text{ClearSky,} & G_t \geq g_2. \end{cases} \quad (9)$$

Values $g_0 = 0$, $g_1 = 150$, and $g_2 = 800 \text{ W m}^{-2}$ are retained only as an illustrative configuration compatible with the current coursework. They are not universal meteorological boundaries. In particular, low irradiance can also be caused by solar position; the four values must therefore be read as generation-context proxies rather than certain weather diagnoses.

Anomaly detection and precedence Let $v_P(t)$, $v_G(t)$, $v_s(t)$, and $v_L(t)$ denote the validity of PV power, tilted irradiance, battery SoC, and aggregate load demand. The control-critical validity predicate is

$$V_t = v_P(t) \wedge v_G(t) \wedge v_s(t) \wedge v_L(t). \quad (10)$$

If V_t is false, no physical anomaly or forecast-driven route is inferred from that observation. A valid irradiance value may still support a descriptive generation context, but the operational validator asserts **SensorMalfunction** and the cycle proceeds directly to safe-mode selection. Within a validated cycle, absence of the three listed

anomaly values is interpreted locally as “none detected”; this is a scoped closed-world policy, not a global OWL assumption. The signed ramp rate, measured in watts per second when Δt is expressed in seconds, is

$$r_t = \frac{P_t - P_{t-1}}{\Delta t}. \quad (11)$$

The anomaly rules are expressed with configurable thresholds: g_{trip} for sufficient irradiance, p_0 for near-zero power, and r_{max} for maximum admissible ramp magnitude.

$$\neg V_t \longrightarrow \text{SensorMalfunction}, \quad (12)$$

$$V_t \wedge G_t \geq g_{\text{trip}} \wedge |P_t| \leq p_0 \longrightarrow \text{InverterTrip}, \quad (13)$$

$$V_t \wedge V_{t-1} \wedge |r_t| > r_{\text{max}} \longrightarrow \text{RampEvent}. \quad (14)$$

No numerical value is inherited for r_{max} because it must be calibrated after invalid values and outliers have been excluded. The controller evaluates sensor validity first. A sensor malfunction therefore cannot be reinterpreted as a physical ramp or inverter trip. If a valid inverter trip also satisfies the ramp predicate, the safe response to the trip takes precedence over ramp smoothing.

Signal observability and structural aliasing For a signal frequency f , sampling frequency f_s , patch size P , and stride S , define a stride-lock ratio and a patch-periodicity ratio

$$\rho_S = \frac{fS}{f_s}, \quad \rho_P = \frac{fP}{f_s}. \quad (15)$$

A stride phase-lock candidate L_S and a within-patch repetition candidate L_P are

$$\begin{aligned} L_S &= [\exists m \in \mathbb{Z}^+ : |\rho_S - m| \leq \varepsilon_S], \\ L_P &= [\exists k \in \mathbb{Z}^+ : |\rho_P - k| \leq \varepsilon_P], \\ C &= (L_S \vee L_P) \wedge \left(0 < f < \frac{f_s}{2}\right), \end{aligned} \quad (16)$$

where ε_S and ε_P are tied to frequency resolution and are never hardcoded frequency bands. Only L_S expresses equality of consecutive patch phases; L_P records whole-cycle repetition inside a patch and is retained as an H3 experimental factor. Their union C is a candidate geometry class, not proof that two domain concepts are indistinguishable.

This is a division of labour, and it bounds what the ontology can conclude on its own. Phase-locking is *derivable*: L_S , L_P and C follow by rule from `pv:signalFrequencyHz`, `pv:samplingFrequencyHz`, `pv:patchSize` and `pv:stride`, all of which the knowledge base already holds, so a reasoner can enumerate the candidate set for any configuration without observing a model. Structural aliasing is *not* derivable and is only recorded: it enters through B below, which is posterior evidence from an experiment external to the ontology, and no rule here can produce it. An agent using this ontology therefore detects phase-locking and reports structural aliasing.

Let q denote the task-relevant information being tested, for example frequency or amplitude. Let $\Delta_{\text{loss}}(q)$ be its predeclared behavioral or representational loss relative to matched controls, δ the maximum practically negligible loss, D the experimental evidence, and $\tau > 1/2$ the posterior confidence threshold. Define confirmed information loss B and confirmed preservation R as

$$B = [\Pr(\Delta_{\text{loss}}(q) > \delta \mid D) \geq \tau], \quad R = [\Pr(\Delta_{\text{loss}}(q) \leq \delta \mid D) \geq \tau]. \quad (17)$$

Because $\tau > 1/2$, B and R cannot both hold for the same assessment. The assessment interface is designed to receive the planned offline evidence already specified in Deliverable 1: a local amplitude-recovery contrast with a heavy-tailed Student- t_4 likelihood and a Gamma regression with log link for non-negative prequential-MDL codelength. The population contrast uses the weakly informative prior $\bar{\beta} \sim \text{Student-}t_4(0, 0.5)$; the Gamma

model uses $\theta_{\text{lock}} \sim \mathcal{N}(0, 0.5)$ and shape $k \sim \text{Gamma}(2, 0.1)$. The stored record identifies the tested (P, S) configuration, matched controls, method and metric, effect estimate, 95% credible interval, and posterior probability of at least 20% attenuation. These are evidence fields, not universal ontology thresholds. Until that planned inference has completed and satisfies the predeclared criterion, no instance may be asserted as **StructurallyAliased**.

The three observability values are assigned as follows:

$$C_O = \begin{cases} \text{StructurallyAliased}, & A, \\ \text{Observable}, & \neg A \wedge R, \\ \text{PartiallyObservable}, & \neg A \wedge \neg R, \end{cases} \quad A = C \wedge B. \quad (18)$$

Observable therefore means that positive evidence shows the tested representation preserves the task-relevant information within tolerance δ . It does not mean global observability of the physical system, nor does it merely mean that aliasing was not found. Candidate geometry may still be **Observable** when the experiment demonstrates preservation. **PartiallyObservable** is the epistemic fallback: preservation has not been certified, while structural aliasing has not been confirmed. It includes a candidate with inconclusive evidence, an unfinished test, a credible interval crossing the decision threshold, conflicting metrics, or no preservation evidence. The term means *partially trusted*, not that a fixed fraction of samples or sensors is visible.

H1 compares each candidate frequency with matched nearby non-candidate controls and predicts higher MDL codelength or lower amplitude recovery; absent or non-localized loss refutes H1. H2 tests whether the loss at a candidate frequency remains approximately stable across sampled phase offsets; systematic phase dependence refutes that hypothesis. H3 tests whether candidate degradation sites move proportionally to $1/S$ or $1/P$ when the corresponding parameter varies. These hypotheses contribute controlled evidence to D , but remain falsifiable experimental factors rather than ontology subclasses or automatic action rules. Sensor validity is modeled separately as an anomaly and never relabels a model representation as structurally aliased.

When an experimental generator provides a phase ϕ in radians, the corresponding temporal offset is converted before storage:

$$o_{\text{samples}} = \frac{\phi f_s}{2\pi f}. \quad (19)$$

Neither ϕ nor o_{samples} is compared directly with stride to infer an anomaly.

The semantic propagation path is explicit:

$$\begin{aligned} \text{pv:TimeSeriesSignal} &\longrightarrow \text{pv:PatchedRepresentation} \\ &\longrightarrow \text{pv:ObservabilityAssessment} \\ &\longrightarrow \text{pv:ControlMode} \longrightarrow \text{pv:PowerRoutingDecision}. \end{aligned} \quad (20)$$

Aliasing can therefore reduce confidence in forecast-derived production, ramp expectations, forecast-driven anomaly recognition, and routing. It does not alter the directly observed values of G_{tilt} , P_{pv} , load demand, or SoC; those values are governed by the separate validity path.

Closed-Loop Control

Mode selection The agent evaluates each synchronized control cycle using the strict priority

$$\text{SafeMode} \succ \text{ConservativeMode} \succ \text{RampOverride} \succ \text{NormalRouting}. \quad (21)$$

Let A_{safe} mean that **SensorMalfunction** or **InverterTrip** is present. The selected mode is

$$M_t = \begin{cases} \text{SafeMode}, & A_{\text{safe}} \vee (C_O = \text{StructurallyAliased}), \\ \text{ConservativeMode}, & C_O = \text{PartiallyObservable}, \\ \text{NormalMode}, & \text{otherwise.} \end{cases} \quad (22)$$

In **SafeMode**, the battery command is Hold and the grid balances the internal load. The PV command is Isolate only for an inverter trip; for sensor malfunction or confirmed structural aliasing it remains connected, but forecast- and ramp-driven battery decisions are suspended. When load demand itself is invalid, Balance carries no invented numerical setpoint. In **ConservativeMode**, the same seven route types remain available with a configurable reduced battery-power cap, normal target/reserve bounds, and low decision confidence. A concurrent ramp cannot activate either extended ramp bound; this implements the precedence in Equation 21.

Routing policy The ordinary policy serves the internal load before charging the battery or exporting energy. Let

$$P_{\text{sur}} = (P_{\text{pv}} - P_{\text{load}})^+, \quad P_{\text{def}} = (P_{\text{load}} - P_{\text{pv}})^+. \quad (23)$$

Here *daylight* abbreviates {ClearSky, PartialCloud, Overcast}. A cycle issues a route only for a strictly positive transfer and may issue several co-mode decisions. All setpoints are clipped to SoC, capacity, and charge/discharge limits.

Figure 1 shows the priority scan and Table 9 states its boundary cases. The conservative branch rejoins the ordinary load-first router under its reduced cap. The ramp branch is available only in **NormalMode**.




Figure 1: Load-first routing decision tree. Elliptical bubbles are ordered gates; rounded bubbles are terminal policies. Safe and conservative decisions precede the normal-mode ramp override, while the ordinary router covers daylight surplus, daylight deficit, and nighttime operation.

Condition	Ordered route set	Constraint
Daylight surplus	PV→Load; PV→Battery; PV→Grid	Serve demand, charge to target, then export the residual.
Daylight deficit	PV→Load; Battery→Load; Grid→Load	Discharge only to reserve; the grid covers the remaining deficit.
Night, $s > s_{\text{reserve}}$	Battery→Load; Grid→Load	Battery serves demand to reserve; the grid supplies the residual.
Night, $s \leq s_{\text{reserve}}$	Grid→Load; if $s < s_{\text{reserve}}$, Grid→Battery	Supply the load first; charge only below reserve and stop at reserve.
Positive ramp with surplus	PV→Load; PV→Battery; PV→Grid	NormalMode may charge beyond target, never beyond maximum SoC.
Negative ramp with valid positive reference	load-serving routes; Battery→Grid	Export only after serving the load, within residual discharge capacity and above reserve.
Negative ramp without usable reference	ordinary surplus/deficit routes	Missing, stale, non-finite, zero, or negative reference disables export catch-up.
ConservativeMode	ordinary load-first routes	Reduced battery cap; target/reserve bounds; no ramp extensions.
SafeMode	Grid→Load when observed	Battery Hold; grid Balance; isolate PV only for inverter trip.

Table 9: Routing matrix for the seven admissible transfer directions.

For a negative ramp, Battery→Grid is admissible only after the internal load has been fully served, the export reference is present, fresh, finite, and positive, measured export is below that reference, the battery remains above reserve, and discharge capacity remains after load support. If any condition fails, the ordinary load-first route applies. At night, Grid→Battery begins only after Grid→Load has been satisfied and stops when reserve is restored. These rules preserve Equation 7 without turning the routing policy into a separate optimisation problem.

References

- [1] Federico Sabbadini. *Chronos-Bolt Patch Sweep*. Model repository: Chronos-Bolt checkpoints retrained from scratch over a grid of patch sizes and strides, one subfolder per geometry. Hugging Face. Aug. 2026. URL: <https://huggingface.co/federicosabbadini/chronos-bolt-patch-sweep> (visited on 08/17/2026).
- [2] *Amazon/Chronos-Bolt-Tiny* · Hugging Face. Dec. 4, 2025. URL: <https://huggingface.co/amazon/chronos-bolt-tiny> (visited on 07/25/2026).
- [3] Elham Tabassi. *Artificial Intelligence Risk Management Framework (AI RMF 1.0)*. NIST AI 100-1. Gaithersburg, MD: National Institute of Standards and Technology (U.S.), Jan. 26, 2023, NIST AI 100-1. DOI: 10.6028/NIST.AI.100-1. URL: <http://nvlpubs.nist.gov/nistpubs/ai/NIST.AI.100-1.pdf> (visited on 06/01/2026).
- [4] Apostol Vassilev et al. *Adversarial Machine Learning: A Taxonomy and Terminology of Attacks and Mitigations*. report NIST AI 100-2e2023. National Institute of Standards and Technology, Jan. 2024. DOI: 10.6028/NIST.AI.100-2e2023. URL: <https://doi.org/10.6028/NIST.AI.100-2e2023> (visited on 08/17/2026).
- [5] Abdul Fatir Ansari et al. *Chronos: Learning the Language of Time Series*. Nov. 4, 2024. DOI: 10.48550/arXiv.2403.07815. arXiv: 2403.07815 [cs.LG]. URL: <http://arxiv.org/abs/2403.07815> (visited on 05/13/2026). Pre-published.